

---

---

# Chord: A Scalable Peer-to-peer Lookup Service for Internet Applications

Ion Stoica, Robert Morris, David Karger, M. Frans Kaashoek, Hari Balakrishnan  
MIT Laboratory for Computer Science

---

---

Presenters: Van Quach & Dale Pedzinski

# Topics

- Introduction
- Related Work
- System Model
- Chord Protocol
- Concurrent Operations and Failures
- Simulation and Experimental Results
- Future Work & Conclusion
- Q & A
- References

# Introduction

- What is Chord Protocol?
- Consistent Hashing
  - Node maintains information about only  $O(\log N)$  nodes
  - Resolves all lookups via  $O(\log N)$  messages
  - Node joins/leave result in no more than  $O(\log^2 N)$  messages
- Main Features of Chord
  - Simplicity
  - Provable Correctness
  - Provable Performance
- Funded by DARPA and the Space and Naval Warfare Systems Center (SPAWAR), San Diego, under contract N66001-00-1-893

# Related Work

- Key/Value Mappings
  - DNS hostname to IP address mapping
  - CAN
- Lookup Operations
  - Freenet peer-to-peer storage system
- Consistent Hashing
  - Ohaha system & Freenet-style query routing
  - Globe System
- Distributed data location protocol
  - Oceanstore
- Lookup Service
  - Napster

# System Model

- Core Features
  - Load Balancing
  - Decentralization
  - Scalability
  - Availability
  - Flexible naming
- Applicable Applications
  - Cooperative Mirroring
  - Time-Shared Storage
  - Distributed Indexes
  - Large-Scale Combinatorial Search

# Chord Protocol

*“Chord Protocol defines how to find the location keys, how new nodes join the system, and how to recover from failure of existing nodes.”*

What is included?

- Consistent Hashing
- Scalable Key Locations
- Node Joins

# Consistent Hashing


*"THEOREM 1. For any set of  $N$  nodes and  $K$  keys, with high probability:*

- 1. Each Node is Responsible for at most  $(1 + e) K/N$  keys*
- 2. When an  $(N + 1)^{st}$  node joins or leaves the network, responsibility for  $O(K/N)$  keys changes hands (and only to or from the joining or leaving node). "*

- hash function assigns each node and key an  $m$ -bit identifier
  - Like SHA-1
- node identifier is node's hashed IP address
- key's identifier is produced by hashing the key □
- Identifiers are ordered in an identifier circle modulo  $2^m - 1$  □
- First node is successor( $k$ ) on the circle

# Consistent Hashing Terms

- Key  $k$  is assigned to the first node with an identifier equal to or in the identifier space □
  - This node = successor node of key □
- Node Joins the Network
  - Keys on  $n$ 's successor move
- Node Leaves the Network
  - Keys move to successor



**Figure 2: An identifier circle consisting of the three nodes 0, 1, and 3. In this example, key 1 is located at node 1, key 2 at node 3, and key 6 at node 0.**



# Scalable Key Locations

*“THEOREM 2. With high probability (or under standard hardness assumptions), the number of nodes that must be contacted to find a successor in an  $N$ -node network is  $O(\log N)$ ”*


- A node only needs to know its successor node on the circle □
- Queries utilize these successor nodes by finding the first node that succeeds the requested identifier □
- It is possible to traverse all  $N$  nodes

# Scalable Key Locations Terms

- **m**=> number of bits in the key/node identifiers  $\square$
- **finger table**=> each node maintains a routing table with, at most, m entries (includes the Chord identifier and the IP address of the relevant node)
  - A node's finger table generally does not contain enough information to determine the successor of a key
- **successor** => first entry of node n's finger table, or the next node on the identifier circle
- **predecessor** => previous node on the identifier circle
- **i<sup>th</sup> finger** =>  $\square$  The i<sup>th</sup> entry of node n's finger table contains the identity of the first node, s, that succeeds n by at least  $2^{i-1}$  on the identifier circle, i.e.,  $s = \text{successor}(n + 2^{i-1})$ , where  $1 \leq i \leq m$  and all arithmetic is modulo  $2^m$

# Search Process

- find\_successor: finds immediate predecessor node of identifier
- Find\_predecessor: moves forward around the Chord circle towards id
- Number of forwards necessary for search:  $O(\log N)$
- Average lookup time is  $1/2 \log N$





# Node Joins

*“THEOREM 3. With high probability, any node joining or leaving an  $N$ -node Chord network will use  $O(\log^2 N)$  messages to re-establish the Chord routing invariants and finger tables. ”*

- Chord preserves two invariants:
  - node's successor is correctly maintained
  - For every key  $k$ , node  $\text{successor}(k)$  is responsible for  $k$
  - (In order for lookups to be fast, it is also desirable for the finger tables to be correct)
- Each node in Chord maintains a predecessor pointer
  - □ contains the Chord identifier and IP address of the immediate predecessor of that node

# Node Joins

- Chord must perform three tasks when a node joins the network: ☐
  - Initialize the fingers and predecessor of node n
  - Update the fingers and predecessors of existing nodes to reflect the addition of n
  - Notify the higher layer software so that it can transfer state (e.g. values) associated with keys that node n is now responsible for




# Initialize the fingers and predecessor of node $n$

- Node  $n$  learns its predecessors by asking  $n'$  to look them up
  - $O(m \log N)$
  - Number of finger entries looked up via utilizing a check on  $i$ th finger entry:  $O(\log N)$
  - Overall time:  $O(\log^2 N)$
- Optimization: a newly joined node  $n$  can ask immediate neighbor for a copy of its finger table and its predecessor
  - Reduces overall time to  $O(\log N)$

# Update the fingers and predecessors

- Node  $n$  will need to be entered into the finger tables of some existing nodes  $\square$
- Node  $n$  will become the  $i$ th finger of node  $p$  iff
  - $p$  precedes  $n$  by at least  $2^{i-1}$   $\square$
  - The  $i$ th finger on node  $p$  succeeds  $n$
- Find predecessor  $p$  of node  $n$ 
  - Check if table needs updates
  - Repeat for predecessor of  $i$
- Number of nodes that need an update:  $O(\log n)$
- Overall time it takes:  $O(\log^2 N)$



# Transfer State

- Move data, associated with each key, over to new node
- Node  $n$  only needs to contact one node to transfer keys to it



# Dynamic Operations and Failures

- Practical issues:
  - a. Nodes join system concurrently
  - b. Nodes fail or leave voluntarily

# Stabilization protocol

Main Goal: Keep nodes' successor pointers up to date

Stabilization scheme:

1. join()
2. stabilize()
3. notify()
4. fix\_fingers()


# join()

- Asks *m* to find the immediate successor of *n*.
- Doesn't make rest of the network aware of *n*.

*n.join(m)*

*predecessor = nil;*

*successor = m.find\_successor(n);*



# stabilize()

- Called periodically to learn about new nodes
- Asks  $n$ 's immediate successor about successor's predecessor  $p$ 
  - Checks whether  $p$  should be  $n$ 's successor instead
  - Also notifies  $n$ 's successor about  $n$ 's existence, so that successor may change its predecessor to  $n$ , if necessary


***$n.stabilize()$***

*$x = \text{successor.predecessor};$*

*if ( $x \in (n, \text{successor})$ )*

*$\text{successor} = x;$*

*$\text{successor.notify}(n);$*



## notify()

- m thinks it might be n's predecessor

*n.notify(m)*

*if (predecessor is nil or  $m \in (\text{predecessor}, n)$ )*

*predecessor = m;*

## fix\_fingers()

- Periodically called to make sure that finger table entries are correct
  - New nodes initialize their finger tables
  - Existing nodes incorporate new nodes into their finger tables

***n.fix\_fingers()***


*next = next + 1 ;*

*if (next > m)*

*next = 1 ;*

*finger[next] = find\_successor(n + 2<sup>next-1</sup>);*

# Stabilization after join



## 1. $n$ joins

- predecessor = nil
- $n$  acquires  $n_s$  as successor
- $n$  notifies  $n_s$  being the new predecessor
- $n_s$  acquires  $n$  as its predecessor

## 2. $n_p$ runs stabilize

- $n_p$  asks  $n_s$  for its predecessor (now  $n$ )
- $n_p$  acquires  $n$  as its successor
- $n_p$  notifies  $n$
- $n$  will acquire  $n_p$  as its predecessor

- all predecessor and successor pointers are now correct**
- fingers still need to be fixed, but old fingers will still work**

# Lookups before stabilization finish


Three behaviors can occur before Chord ring stable:

1. All finger table entries are reasonably current  $\Rightarrow$  take  $O(\log N)$  steps
2. Successor pointers are correct, but fingers inaccurate  $\Rightarrow$  correct lookups, but slower
3. Incorrect successor, or keys not yet migrated to newly joined nodes  $\Rightarrow$  lookups may fail, can pause and retry.



# Failure Recovery

- Main goal: maintain correct successor pointers
- How:
  - Each node maintains a successor list of  $r$  nearest successors on the ring
  - If node  $n$  notices its successor has failed, it replaces the failed node with the 1st live entry
  - *stabilize* will correct finger table entries and successor-list entries pointing to failed node



# Voluntary Node Departures


- Can be treated as node failures with enhancements
  - Transfer all keys to its successor
  - Notify its predecessor and successor

## Chord - Fact

- Every node is responsible for about  $K/N$  keys
- When a node joins or leaves an  $N$ -node network, only  $O(K/N)$  keys change hands (and only to and from joining or leaving node)
- Lookups need  $O(\log N)$  messages
- Stabilization for node leaving / joining need  $O(\log^2 N)$  message

# Experimental Results

1. Latency grows slowly with total number of nodes
2. Chord is robust for multiple node failures



# Q & A

# References

[https://pdos.csail.mit.edu/papers/chord:sigcomm01/chord\\_sigcomm.pdf](https://pdos.csail.mit.edu/papers/chord:sigcomm01/chord_sigcomm.pdf)