

Spanner: Google's Globally-Distributed Database

Cheng Shiyang
Samuel Cherinet



Introduction

What is Spanner?

- Highly Scalable
- Semi-Relational

Features

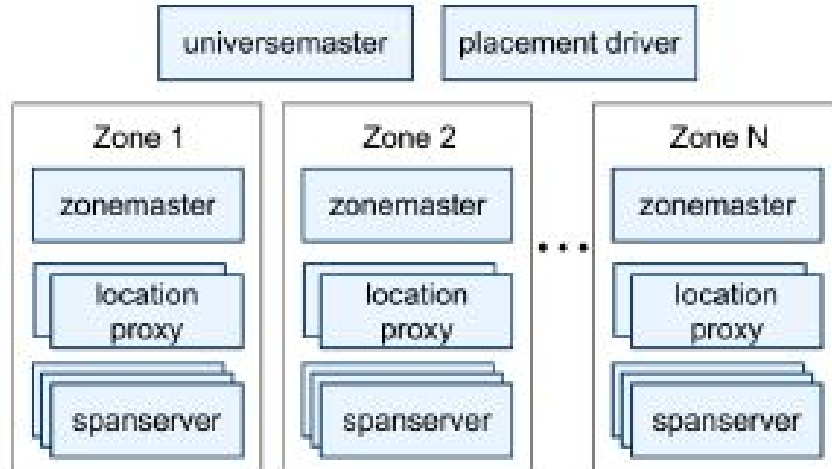
- Replication Configuration
- Externally consistent read/write
- Globally meaningful commit timestamps
- SQL

True Time API

Cloud Spanner: The best of the relational and NoSQL worlds

| | CLOUD SPANNER | TRADITIONAL RELATIONAL | TRADITIONAL NON-RELATIONAL |
|--------------|---------------------|------------------------|----------------------------|
| Schema | ✓ Yes | ✓ Yes | ✗ No |
| SQL | ✓ Yes | ✓ Yes | ✗ No |
| Consistency | ✓ Strong | ✓ Strong | ✗ Eventual |
| Availability | ✓ High | ✗ Failover | ✓ High |
| Scalability | ✓ Horizontal | ✗ Vertical | ✓ Horizontal |
| Replication | ✓ Automatic | ↻ Configurable | ↻ Configurable |

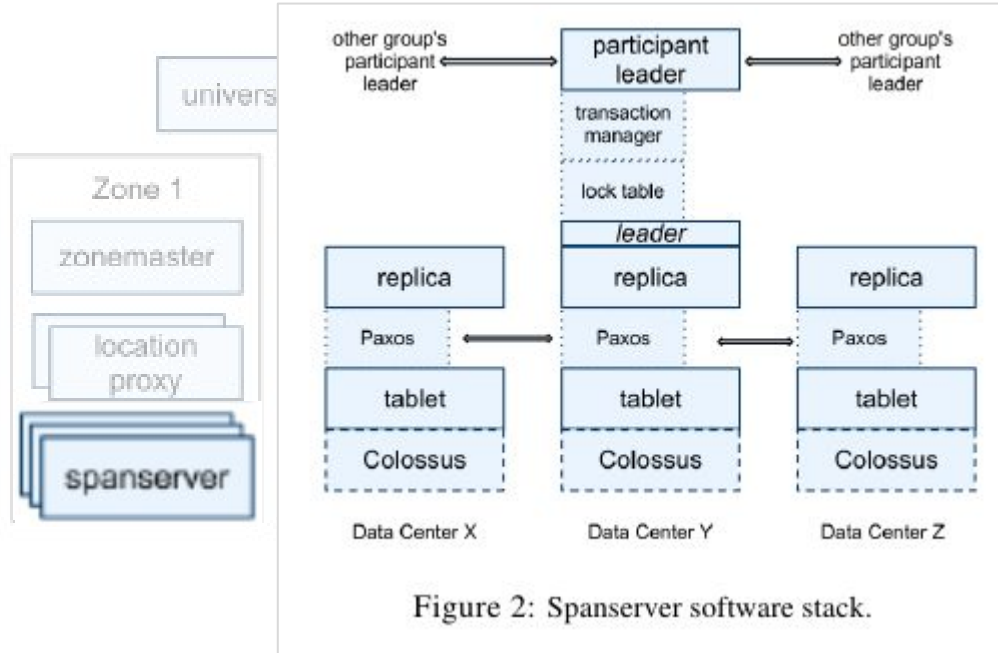
Implementation : Universe



- ❑ Spanserver
- ❑ Zonemaster
- ❑ Location Proxy

Figure 1: Spanner server organization.

Spanserver



- ❑ Tablet
- ❑ Distributed Transaction
- ❑ Replication
- ❑ Colossus

Directories




Figure 3: Directories are the unit of data movement between Paxos groups.

- ❑ Prefixing for locality
- ❑ Movedir
- ❑ Administration
 - ❑ number/type of replicas
 - ❑ Geographic placement

Data Model

- ❑ Schematized semi-relational tables :
 - ❑ Rows must have name (one or more primary key columns)
- ❑ Query Language
 - ❑ Similar to SQL but has support for protocol-buffer-values fields
- ❑ Strongly Typed
- ❑ Layered on top of the directory-bucketed key-value mappings supported by the implementation

Example Physical Layout of Data

| SingerId | First Name | AlbumId | SongId | Album Title | Song Title |
|----------|------------|---------|--------|---------------------------|---------------------------|
| 1 | "Marc" | 1 | 1 | "Total Junk" | |
| 1 | "Marc" | 2 | 1 | "Go, Go, Go" | |
| 1 | "Marc" | 1 | 2 | | "42" |
| 1 | "Marc" | 2 | 2 | | "Nothing Is The Same" |
| 2 | "Catalina" | 2 | 1 | "Green" | |
| 2 | "Catalina" | 1 | 1 | | "Let's Get Back Together" |
| 2 | "Catalina" | 1 | 2 | | "Starting Again" |
| 2 | "Catalina" | 1 | 3 | | "I Knew You Were Magic" |
| 2 | "Catalina" | 2 | 2 | "Forever Hold Your Peace" | |
| 2 | "Catalina" | 2 | 3 | "Terrified" | |
| 2 | "Catalina" | 3 | 1 | | "Fight Story" |

Predecessor/Other Data Models

- **Big Table**
- **Megastore**
 - Over 300 application use it in google
 - Relative low performance
 - Synchronous replication


True Time

True Time APIs:

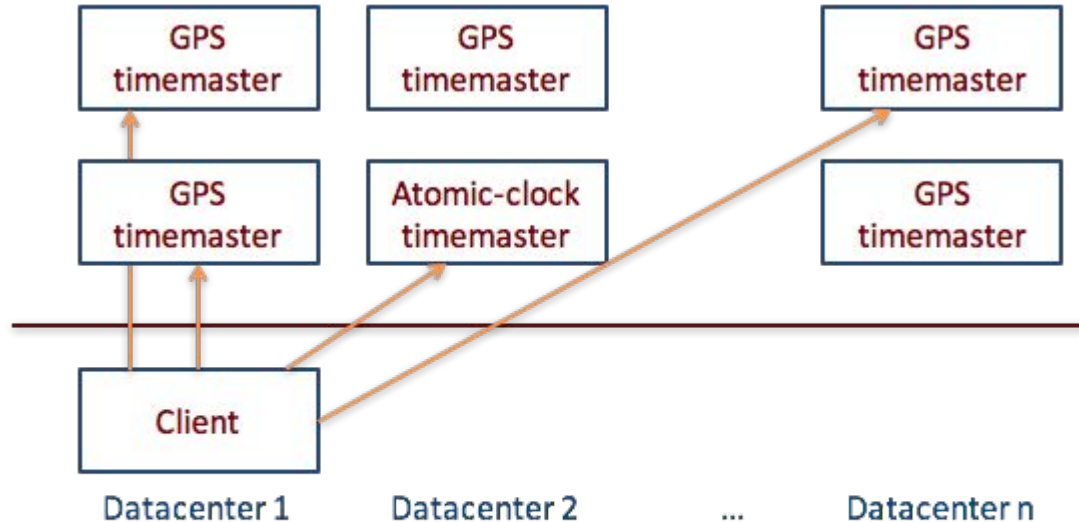
TT.now() // [earliest, latest]

TT.after(t)

TT.before(t)



True Time Architecture



Compute reference $[\text{earliest}, \text{latest}] = \text{now} \pm \epsilon$

Marzullo's Algorithm to detect and reject liars.

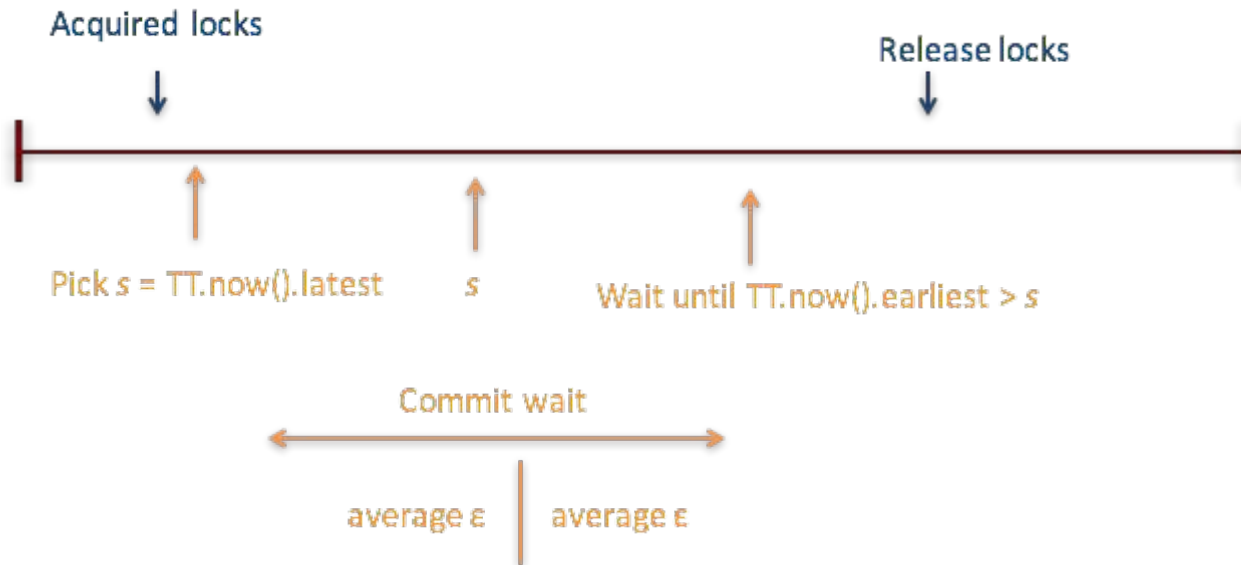
External Consistency

Transaction: T1, T2

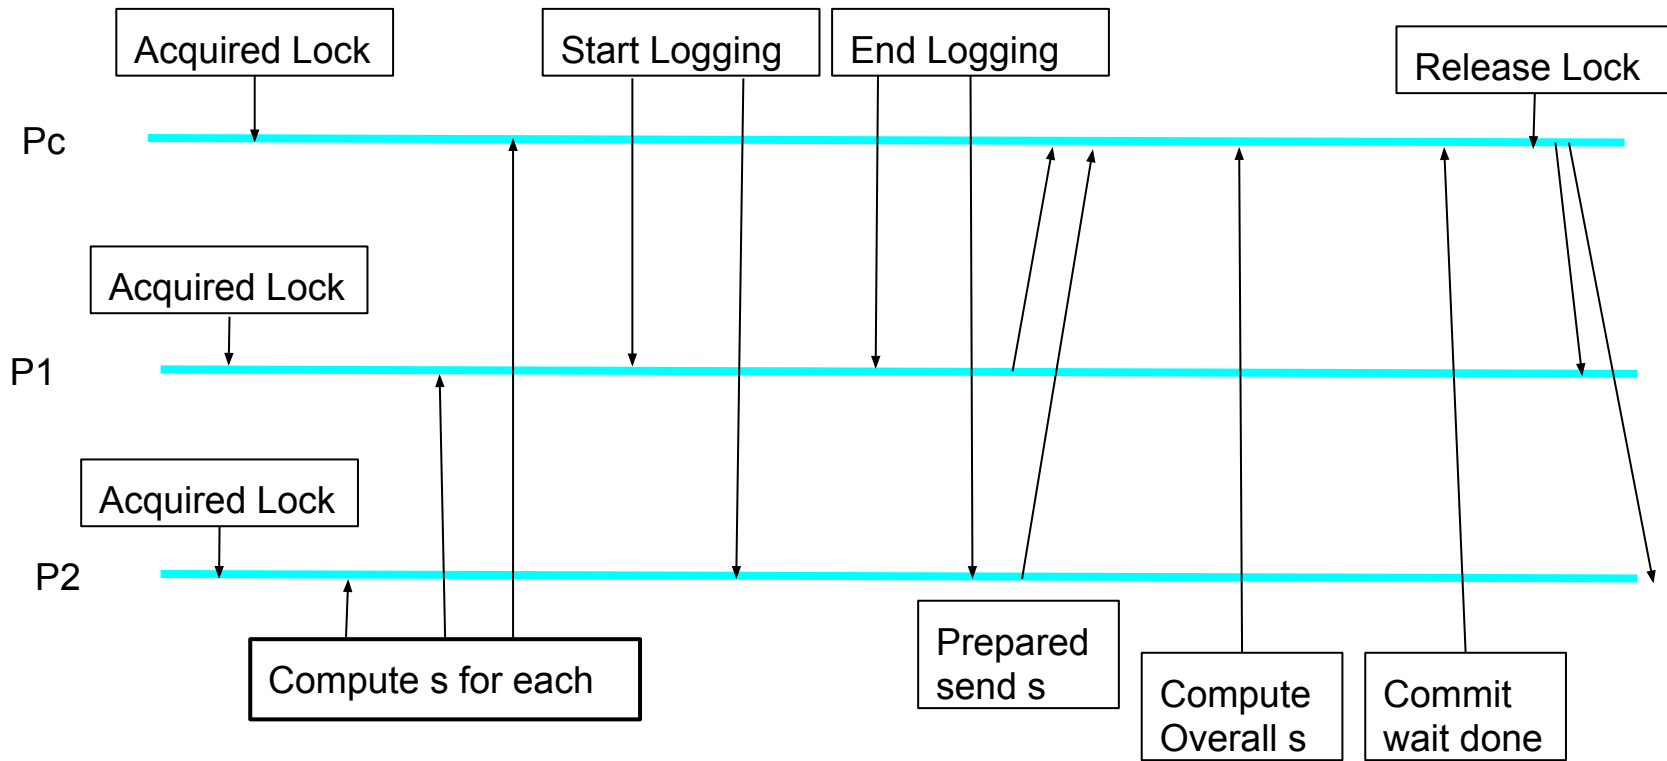
If T1 commits before T2's

$T1's\ ts < T2's\ ts$

Timestamps and TrueTime



Commit Wait and 2-Phase Commit



Read-Write Transaction

- Writes occur in a transaction are buffered at the client until commit
- Read transaction will acquires read locks and then read most recent data
- After read transaction client will apply two phase-commit
- A non-coordinator-participant leader first acquires write locks, choose a timestamp then logs the prepare record
- Coordinator participant first acquire the lock, skip the prepare phase and log the record
- Before allowing any coordinator commit, it will wait for $TT.after(s)$
- Any coordinator commit and release the lock

Read-Only Transaction

- Define LastTS() to be the timestamp of the last committed
- In single-Paxos group, If there is no prepare transactions, the assignment $s = \text{LastTS}()$
- In multi-Paxos group, assign $s = \text{TT.now().latest}$


Schema-Change Transaction

- Assign a timestamp in future, which registered in the prepare timestamp at time t
- Reads and writes may proceed if their timestamp precede t
- Reads and writes must block behind the schema-change transaction if their timestamp are after t


Evaluation - Two phase commit Scalability

| participants | latency (ms) | |
|--------------|------------------|------------------|
| | mean | 99th percentile |
| 1 | 17.0 \pm 1.4 | 75.0 \pm 34.9 |
| 2 | 24.5 \pm 2.5 | 87.6 \pm 35.9 |
| 5 | 31.5 \pm 6.2 | 104.5 \pm 52.2 |
| 10 | 30.0 \pm 3.7 | 95.6 \pm 25.4 |
| 25 | 35.5 \pm 5.6 | 100.4 \pm 42.7 |
| 50 | 42.7 \pm 4.1 | 93.7 \pm 22.9 |
| 100 | 71.4 \pm 7.6 | 131.2 \pm 17.6 |
| 200 | 150.5 \pm 11.0 | 320.3 \pm 35.1 |

Evaluation - Availability



Evaluation - True Time



Question?